

Collect distributed application logging

using Fluentd (EFK stack)

Marco Pas

Philips Lighting

Software geek, hands on
Developer/Architect/DevOps Engineer

@marcopas

Some stuff about me...

- Mostly doing cloud related stuff
 - Java, Groovy, Scala, Spring Boot, IOT, AWS, Terraform, Infrastructure
- Enjoying the good things
- Chef leuke dingen doen == “trying out cool and new stuff”
- Currently involved in a big IOT project
- Wannabe chef, movie & Netflix addict


Agenda

- Logging
- Distributed Logging
- Fluentd Overview including demo's:
 - Run Fluentd
 - Capture input from docker container
 - Capture HTTP access logs
 - Capture HTTP access logs and store in MongoDB
 - Capture HTTP access logs and store in EFK stack
 - Capture SpringBoot logs and store in EFK stack including in_tail
 - HA Setup

Logging



- Providing useful information, seems hard!
- Common Log Formats
 - W3C, Common Log Format, Combined Log Format
 - used for:
 - Proxy & Web Servers
- Agree upon Application Log Formats
 - Do not forget -> Log levels!
- Data security
 - Do not log passwords or privacy related data



Some seriously useful log message :)

- “No need to log, we know what is happening”
- “Something happened not sure what”
- “Empty log message”
- “Lots of sh*t happing”
- “It works b****”
- “How did we end up here?”
- “Okay i am getting tired of this error message”
- “Does this work?”
- “We hit a bug, still figuring out what”
- “Call 911 we have a problem”

Logging considerations

- Logging means more code
- Logging is not free
- Consider feedback to the UI instead of logging
- **The more you log, the less you can find**
- Consider to log only the most evil scenarios (log exceptions)
- Agree on levels like FATAL, ERROR, WARN, DEBUG, INFO, TRACE ...

LOG ALL THE THINGS





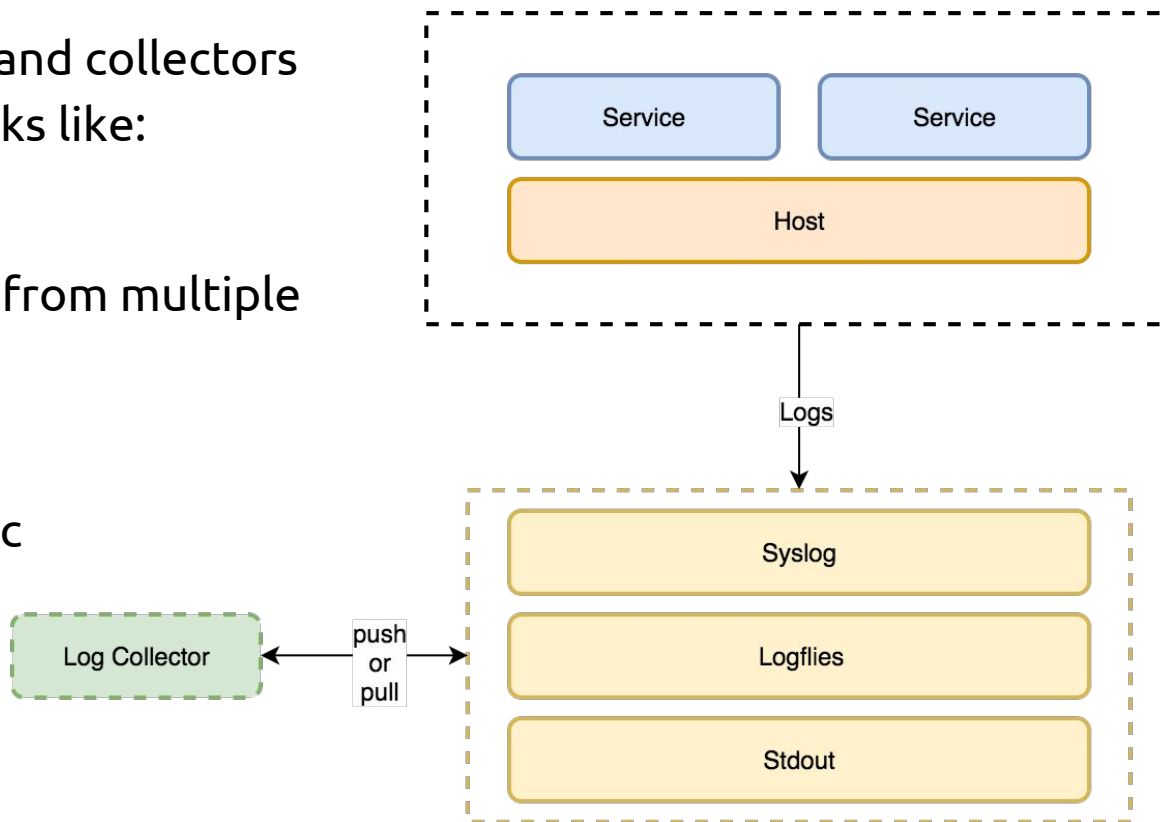
- Syslog / Syslog-ng
- Files -> multiple places (/var/log)
 - Near realtime replication to remote destinations
- Stdout
 - Normally goes to /dev/null



In container based environments logging to “Stdout” has the preference



- Specialized transporters and collectors available using frameworks like:
 - Logstash, Flume, Fluentd
- Accumulate data coming from multiple hosts / services
 - Multiple input sources
- Optimized network traffic
 - Pull / Push






- Where should it be stored?
 - Short vs Long term
 - Associated costs
 - Speed of data ingestion & retrieval
 - Data access policies (who needs access)

- Example storage options:


- S3, Glacier, Tape backup
- HDFS, Cassandra, MongoDB or ElasticSearch





- Batch processing of log data
 - HDFS, Hive, PIG → MapReduce Jobs

- UI based Analyses
 - Kibana, GrayLog2





- Based on patterns or “calculated” metrics → send out events
 - Trigger alert and send notifications
- Logging != Monitoring
 - Logging -> recording to diagnose a system

```
127.0.0.1 - frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326
```

- Monitoring -> observation, checking and recording

```
http_requests_total{method="post",code="200"} 1027 1395066363000
```

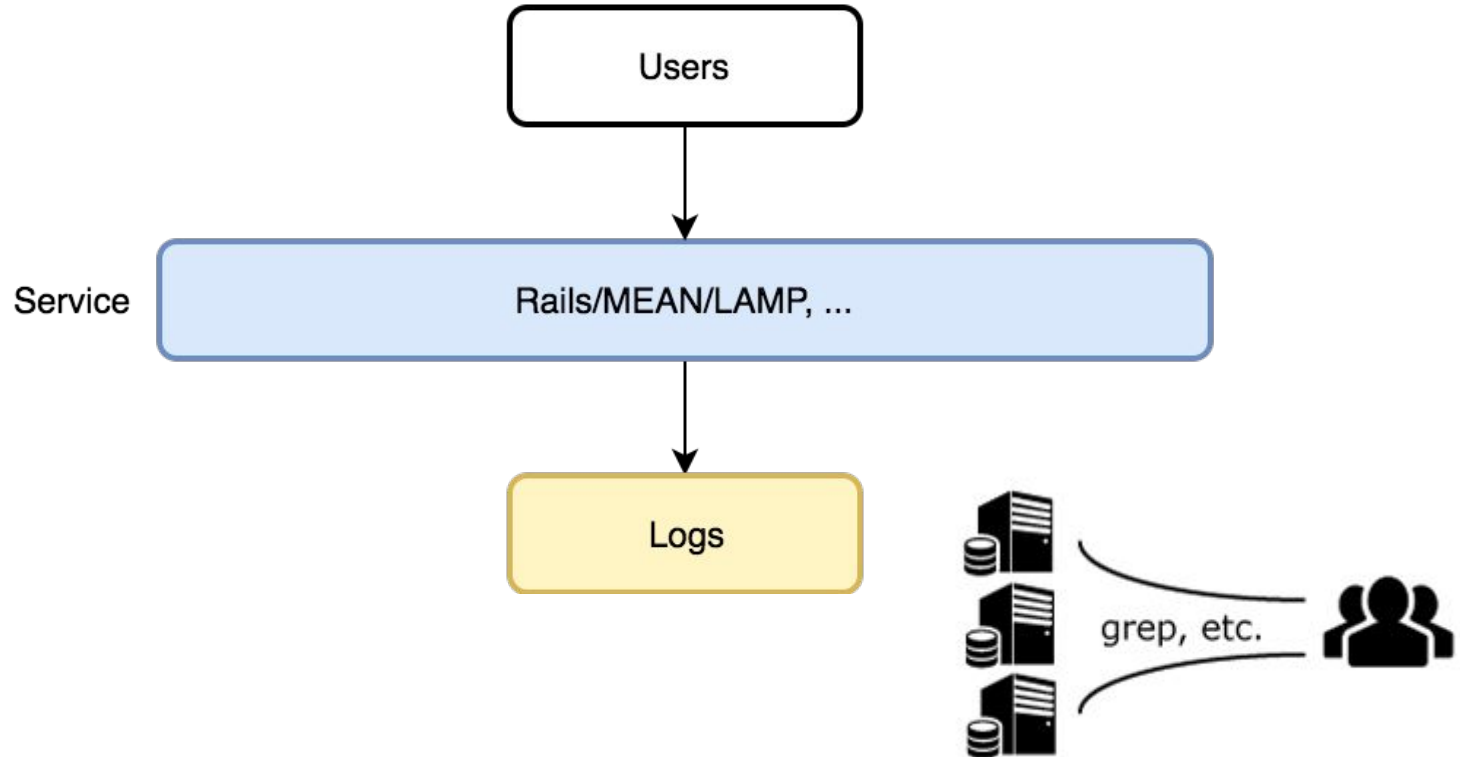
*“In a containerized world,
we must think differently
about logging.”*


Label data at the source


Push data and parse it as soon as possible

Distributed Logging

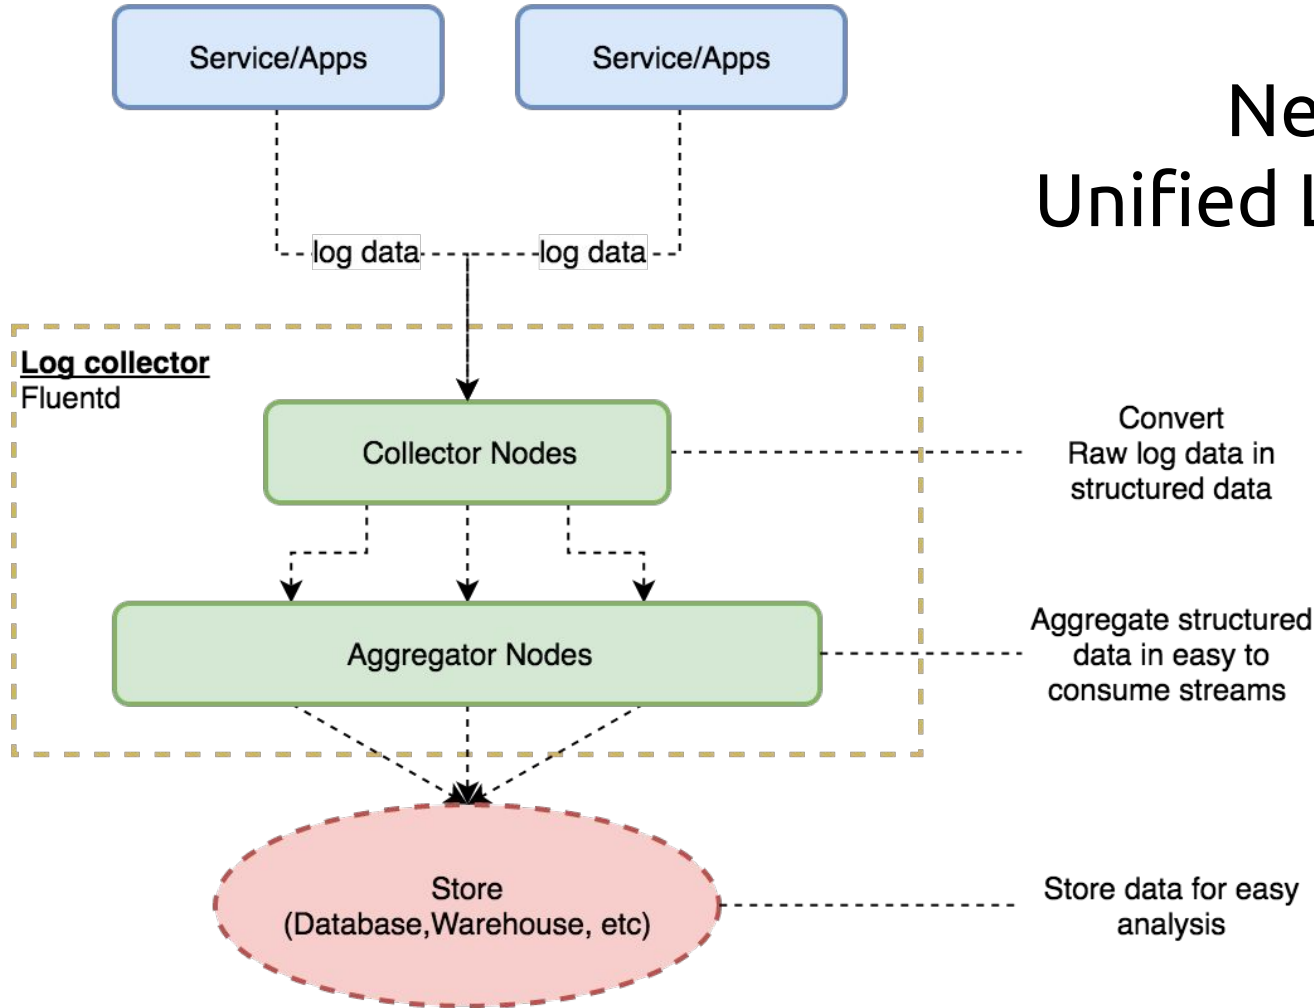
Logging







Need for a Unified Logging Layer



Fluentd Overview

Access logs

Apache

App logs

Frontend
Backend

System logs

syslogd

Databases



fluentd



filter / buffer / routing

Alerting

Nagios

Analysis


MongoDB
MySQL
Hadoop

Archiving

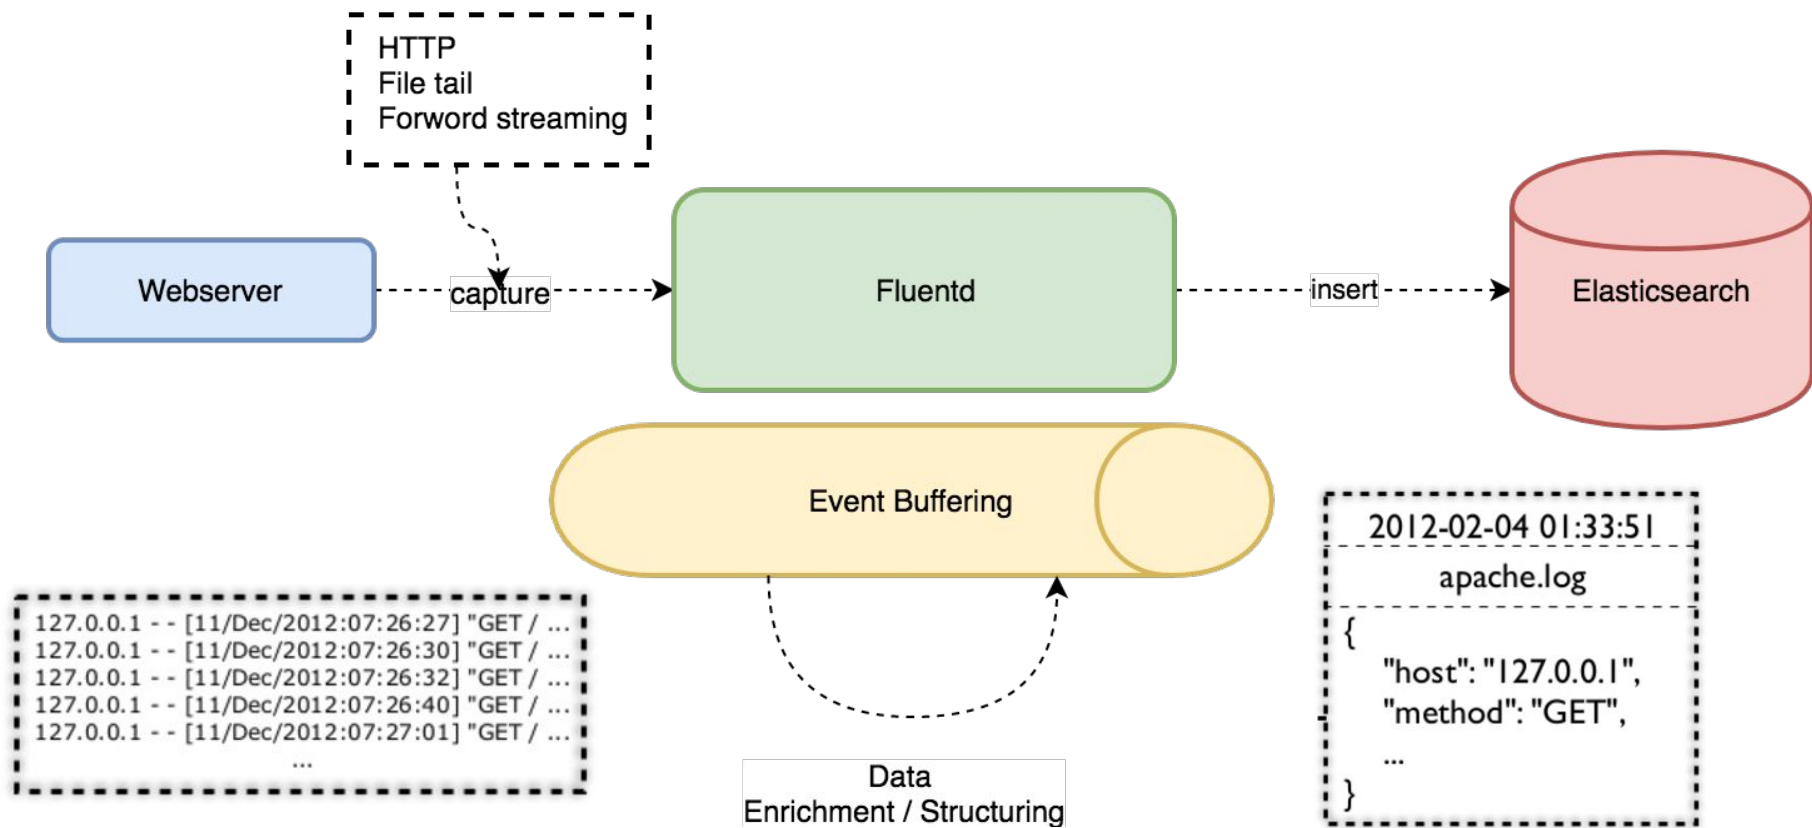
Amazon S3

Fluentd

- Open source log collector written in Ruby
- Reliable, scalable and easy to extend
 - Pluggable architecture
 - Rubygem ecosystem for plugins
- Reliable log forwarding

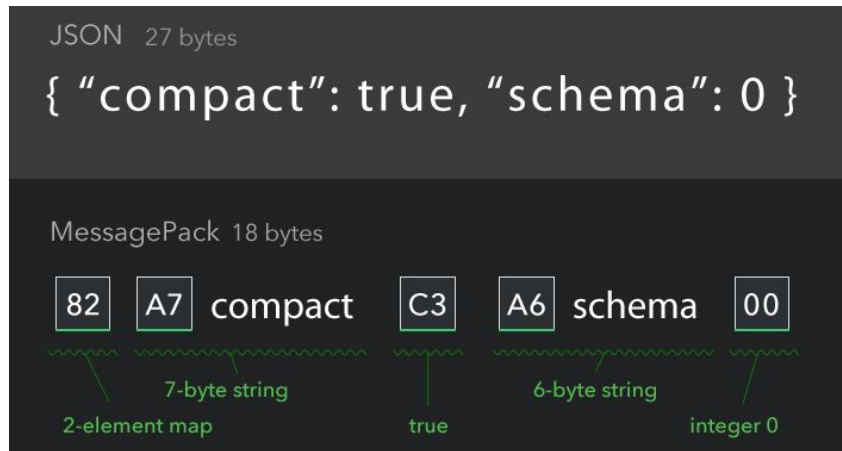


Example



Event structure

- Tag
 - Where an event comes from, used for message routing
- Time
 - When an event happens, Epoch time
 - Parsed time coming from the datasource
- Record
 - Actual log content being a JSON object
 - Internally MessagePack



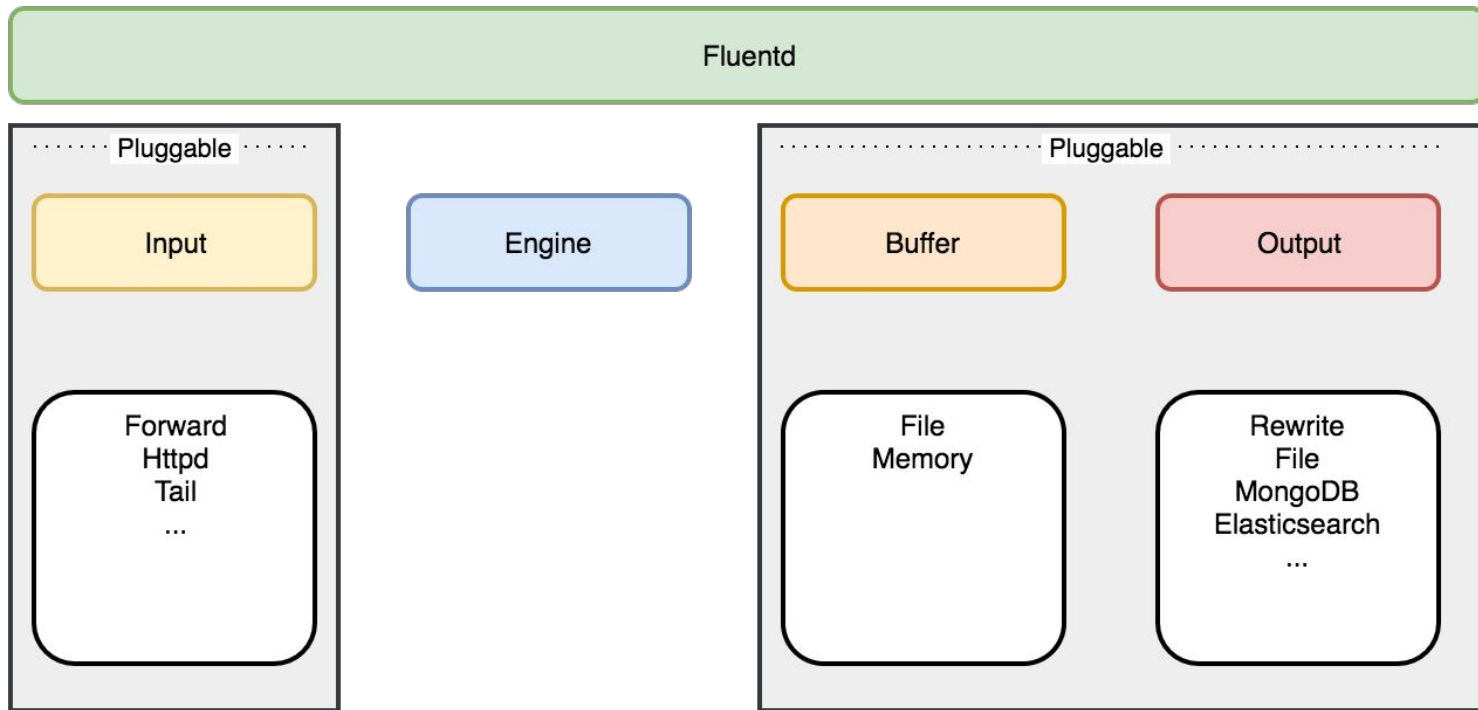
Event example

```
192.168.0.1 - - [28/Feb/2013:12:00:00 +0900] "GET / HTTP/1.1" 200 777
```



```
tag:: apache.access # set by configuration  
time: 1362020400    # 28/Feb/2013:12:00:00 +0900  
record: {"user":"-","method":"GET","code":200,"size":777,"host":"192.168.0.1","path":"/"}
```


Pluggable Architecture



<http://www.fluentd.org/plugins>

Configuration

- Driven by a simple text based configuration file
 - fluent.conf

```
<source><source/>
```

→ Tell where the data comes from (input)

```
<match></match>
```


→ Tell fluentd what to do (output)

```
<filter></filter>
```

→ Event processing pipeline

```
<label></label>
```

→ Groups filter and output for internal routing



source -> filter 1 -> ... -> filter N -> output

```
# receive events via HTTP
```

```
<source>
```

```
  @type http
```

```
  port 9880
```

```
</source>
```

```
# read logs from a file
```

```
<source>
```

```
  @type tail
```

```
  path /var/log/httpd.log
```

```
  format apache
```

```
  tag apache.access
```

```
</source>
```

```
# save alerts to a file
```

```
<match alert.**>
```

```
  @type file
```

```
  path /var/log/fluent/alerts
```

```
</match>
```

```
# save access logs to MongoDB
```

```
<match apache.access>
```

```
  @type mongo
```

```
  database apache
```

```
  collection log
```

```
</match>
```

```
# forward other logs to servers
```

```
<match **>
```

```
  type forward
```

```
<server>
```

```
  host 192.168.0.11
```

```
  weight 20
```

```
</server>
```

```
<server>
```

```
  host 192.168.0.12
```

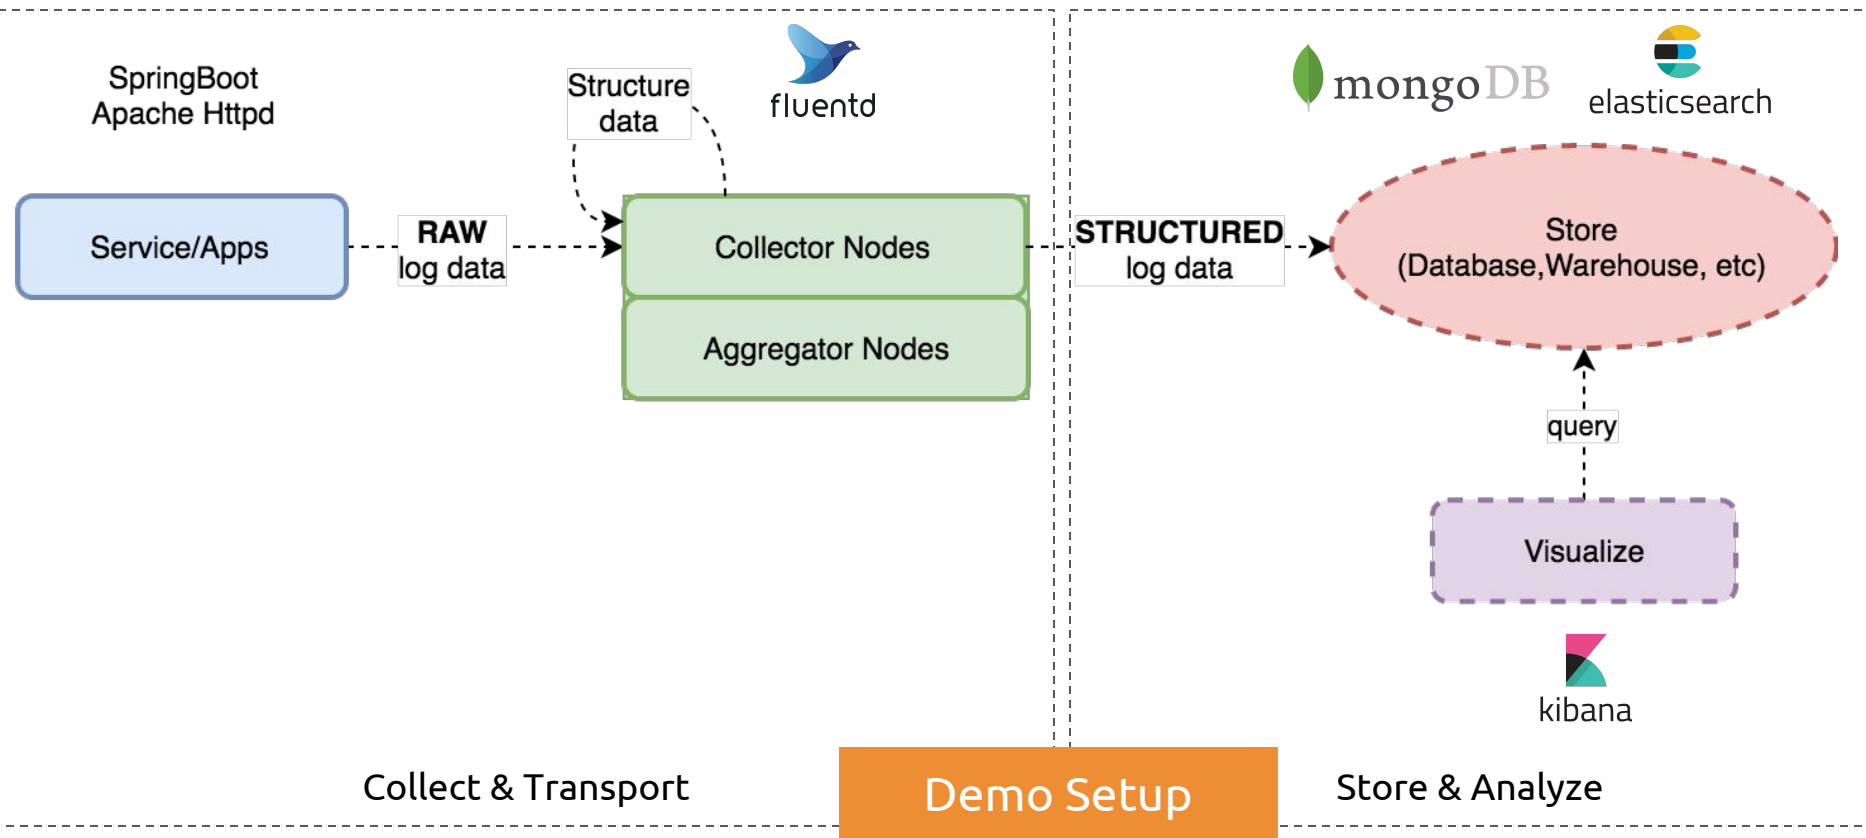
```
  weight 60
```

```
</server>
```

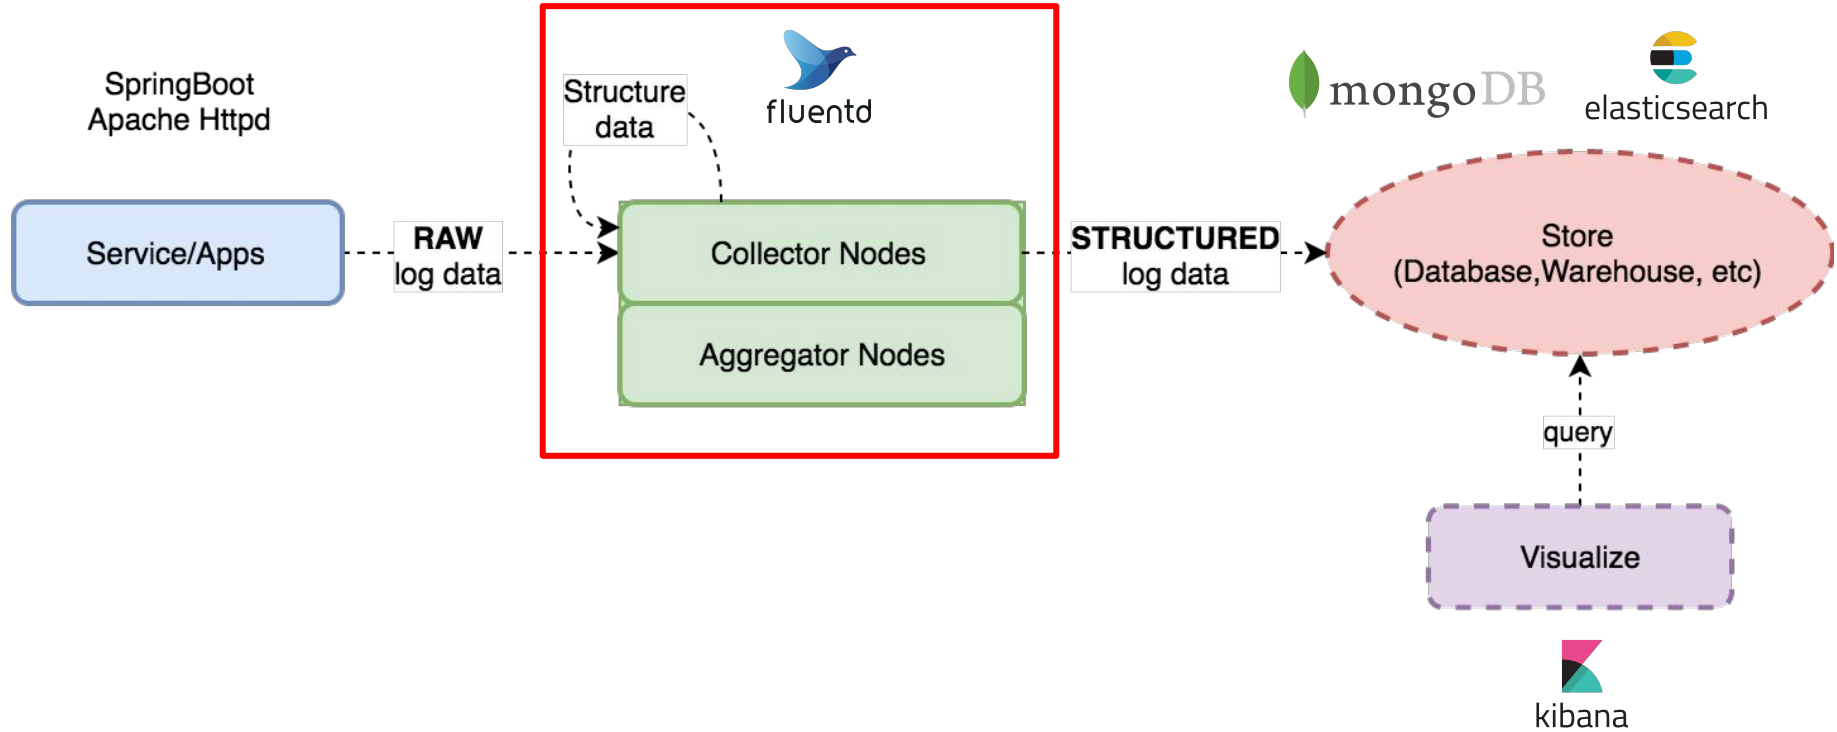
```
</match>
```

```
# add a field to an event
<filter myapp.access>
  @type record_transformer
  <record>
    host_param "#{Socket.gethostname}"
  </record>
</filter>
```

```
# grouping and internal routing
<source>
  @type forward
  port 24224
  bind 0.0.0.0
  @label @SYSTEM
</source>
<label @SYSTEM>
  <filter var.log.middleware.**>
    @type grep
    # ...
  </filter>
  <match **>
    @type s3
    # ...
  </match>
</label>
```



Demo: Run Fluentd



```
# file: docker-compose.yml
```

```
version: '2'
```

```
services:
```

```
  fluentd:
```

```
    container_name: fluentd
```

```
    image: fluentd-demo
```

→ Docker image used for fluentd (container the plugins)

```
    volumes:
```

```
      - $PWD/./fluentd/etc
```

→ Mounting local filesystem that contains the config file

```
    ports:
```

```
      - "24220:24220"
```

→ portmapping 24220 on host to 24220 in Docker container

```
# file: fluent.conf
```

```
# monitoring agent:
```

```
#   check http://localhost:24220/api/plugins.json for healthcheck
```

```
<source>
```

```
  @type monitor_agent
```

```
  port 24220
```

→ Run the **monitor agent** on port 24220

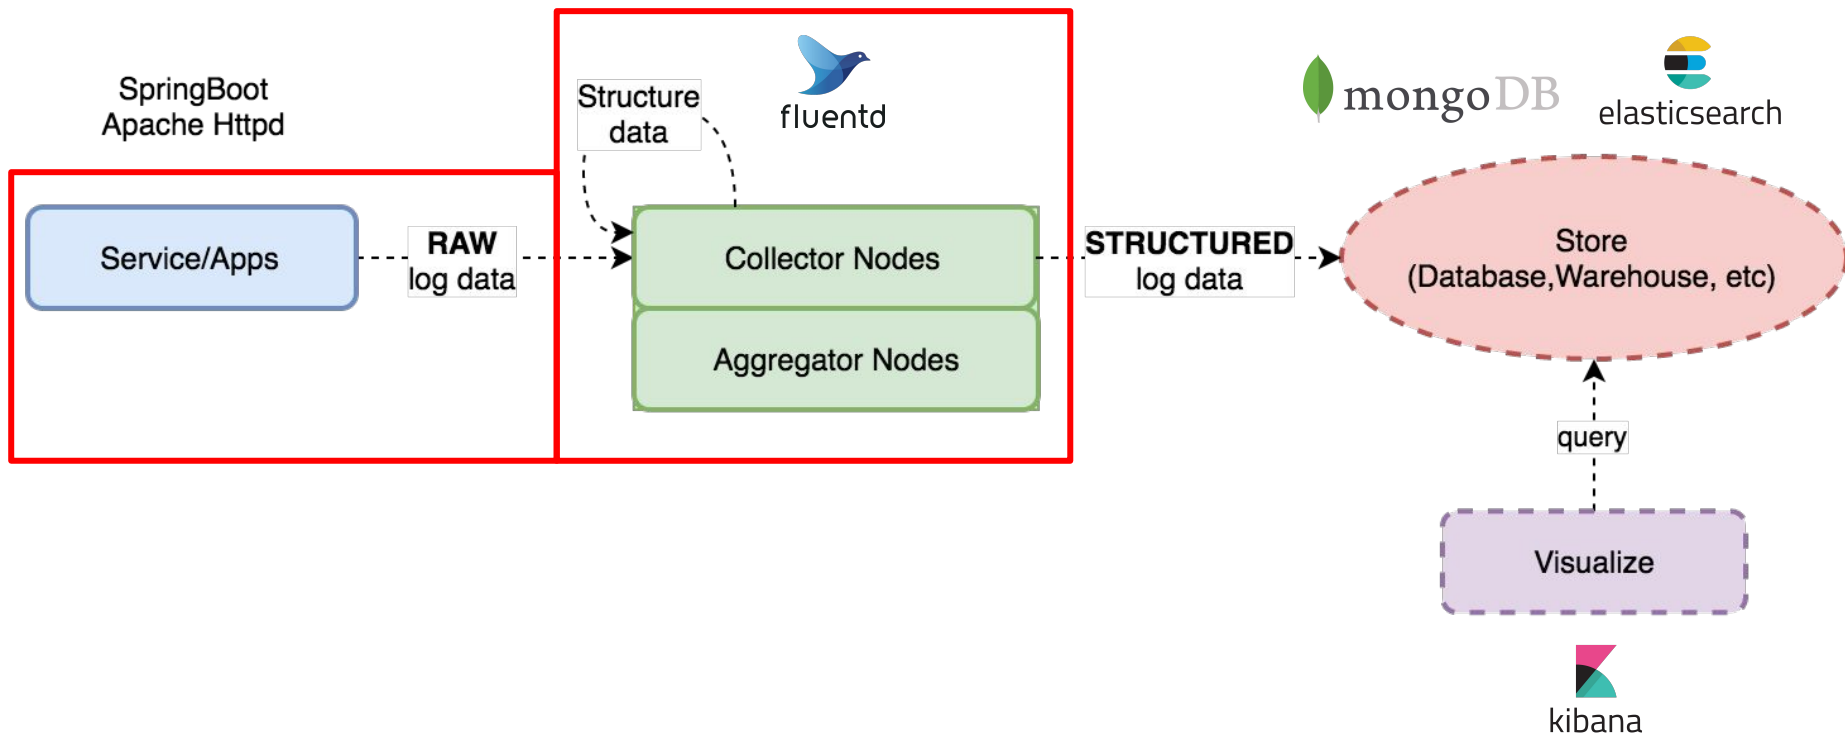
```
  bind 0.0.0.0
```

→ Bind to all network interfaces

```
</source>
```


Demo

Demo: Capture input from Docker container



```
# file: docker-compose.yml
```

```
version: '2'
```

```
services:
```

```
  fluentd:
```

```
    container_name: fluentd
```

```
    # code intentionally omitted
```

```
  echo:
```

```
    container_name: echo
```

```
    image: debian
```

```
    command: bash -c 'for((i=1;i<=1000;i+=1)); do echo -e "Welcome ${i} times"; sleep 2; done;'
```

```
    links:
```

```
      - fluentd
```

```
    logging:
```

```
      driver: "fluentd" → Use the fluentd logging driver
```

```
      options:
```

```
        fluentd-address: localhost:24224 → Where can we find fluentd?
```

```
        tag: echo → Tag used for event routing
```

```
# file: fluent.conf
```

```
# input forward plugin
```

```
<source>
```

```
  @type forward
```

→ Bind to all network interfaces

```
  port 24224
```

→ Run the `in_forward` plugin on port 24220

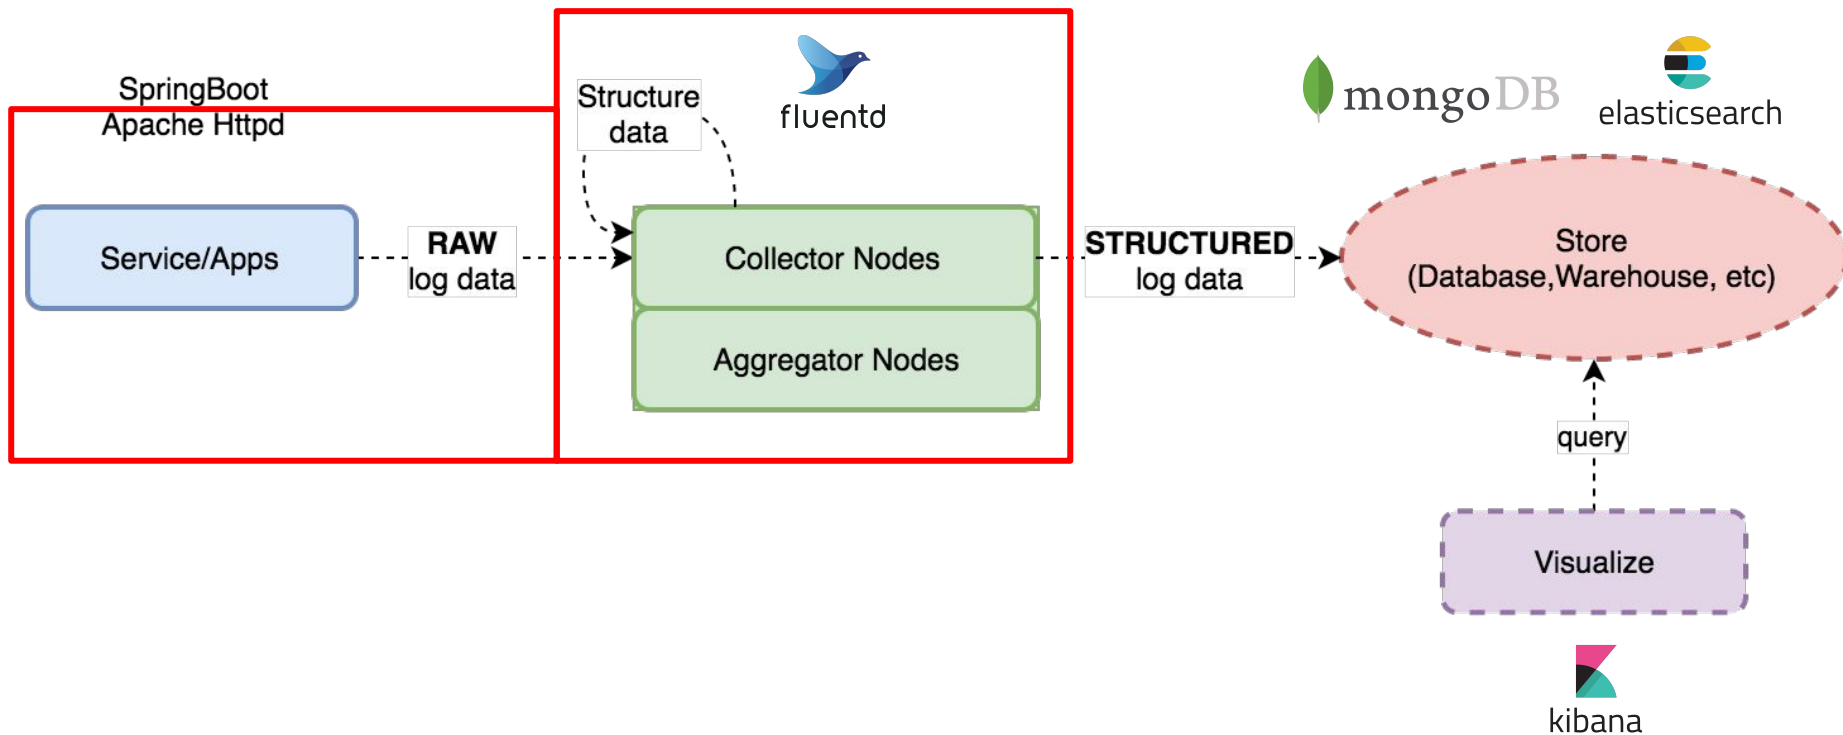
```
  bind 0.0.0.0
```

→ Bind to all network interfaces

```
</source>
```

Demo

Demo: Capture HTTP Access Logs



```
# file: docker-compose.yml
```

```
version: '2'
```

```
services:
```

```
  fluentd:
```

```
    container_name: fluentd
```

```
    # code intentionally omitted
```

```
  httpd:
```

```
    container_name: httpd
```

```
    image: httpd-demo
```

```
    ports:
```

```
      - "80:80"
```

→ Run our Http server on port 80 serving "/"

```
    links:
```

```
      - fluentd
```

```
    logging:
```

```
      driver: "fluentd"
```

→ Use the fluentd logging driver

```
      options:
```

```
        fluentd-address: localhost:24224
```

→ Where can we find fluentd?

```
        tag: httpd.access
```

→ Tag used for event routing

You get the idea :)

```
# file: fluent.conf
# input forward plugin
```

```
<source>
```

```
  @type forward
```

→ Bind to all network interfaces

```
  port 24224
```

→ Run the `in_forward` plugin on port 24220

```
  bind 0.0.0.0
```

→ Bind to all network interfaces

```
</source>
```

```
# filter httpd access logs
```

```
<filter httpd.access>
```

→ Notice the filter tag! `*`, `*.*`, `**`, `{a.b,a.*,a.*.b}`, ...

```
  @type parser
```

→ Parse the data and create fields using the regex pattern

```
  format /^some regex pattern$/
```

```
  # code intentionally omitted
```

```
</filter>
```

```
# match all and print
```

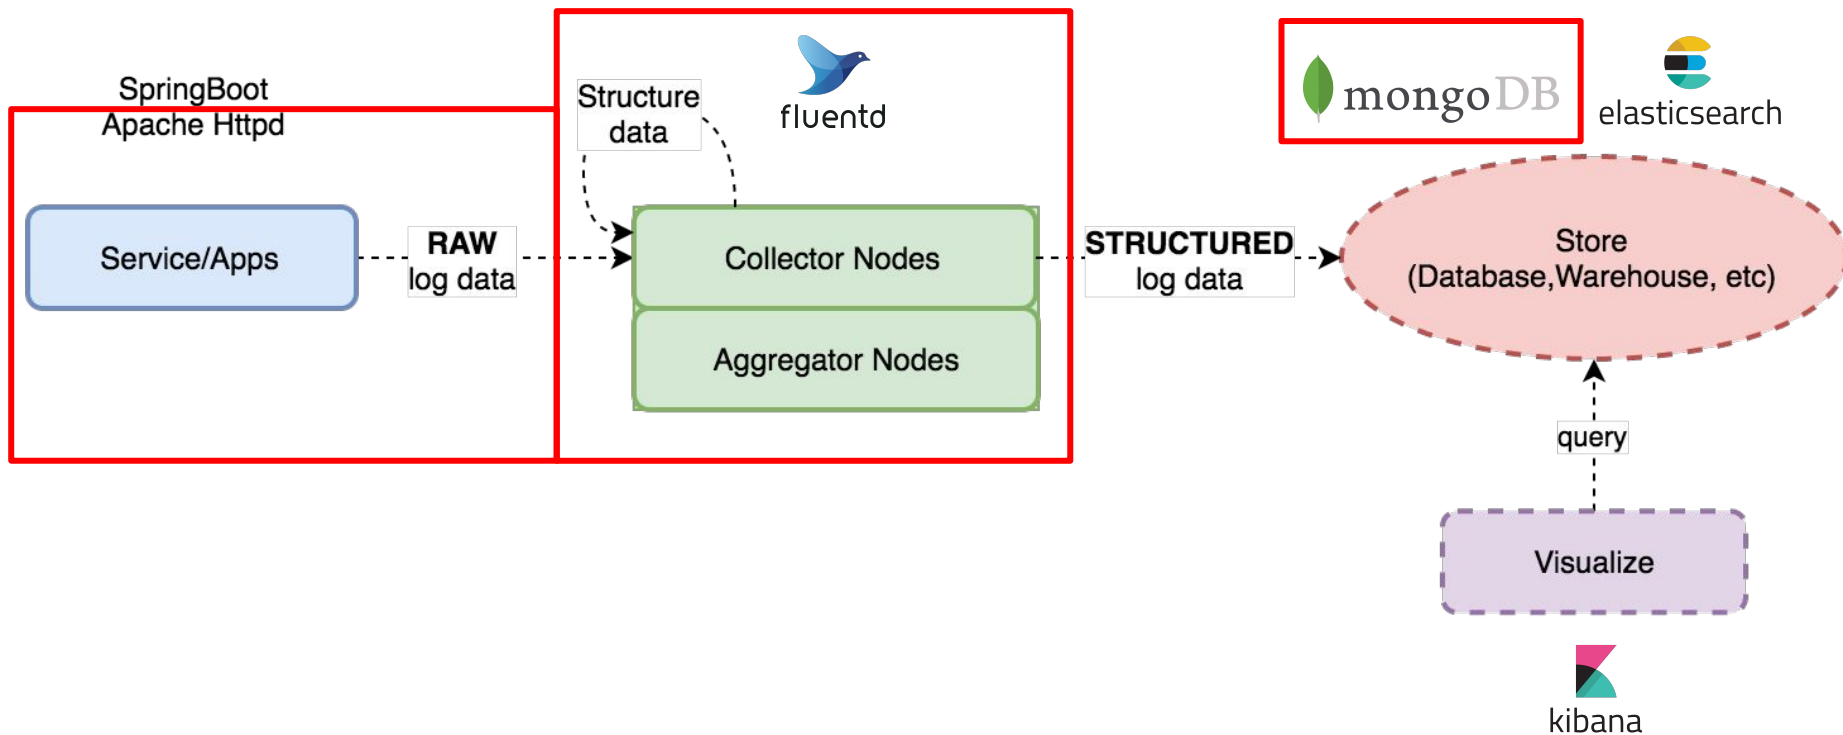
```
<match **>
```

```
  @type stdout
```

```
</match>
```


Demo

Demo: Capture HTTP Access Logs -> MongoDB



```
# file: fluent.conf  
# code intentionally omitted
```

```
<match httpd.access>
```

```
  @type copy                                → Copy to multiple destinations
```

```
  <store>
```

```
    @type stdout                            → Console output
```

```
  </store>
```

```
  <store>
```

```
    @type mongo                            → MongoDB output
```

```
    host mongodb
```

```
    port 27017
```

```
    database fluentd
```

```
    collection test
```

```
    flush_interval 5s
```

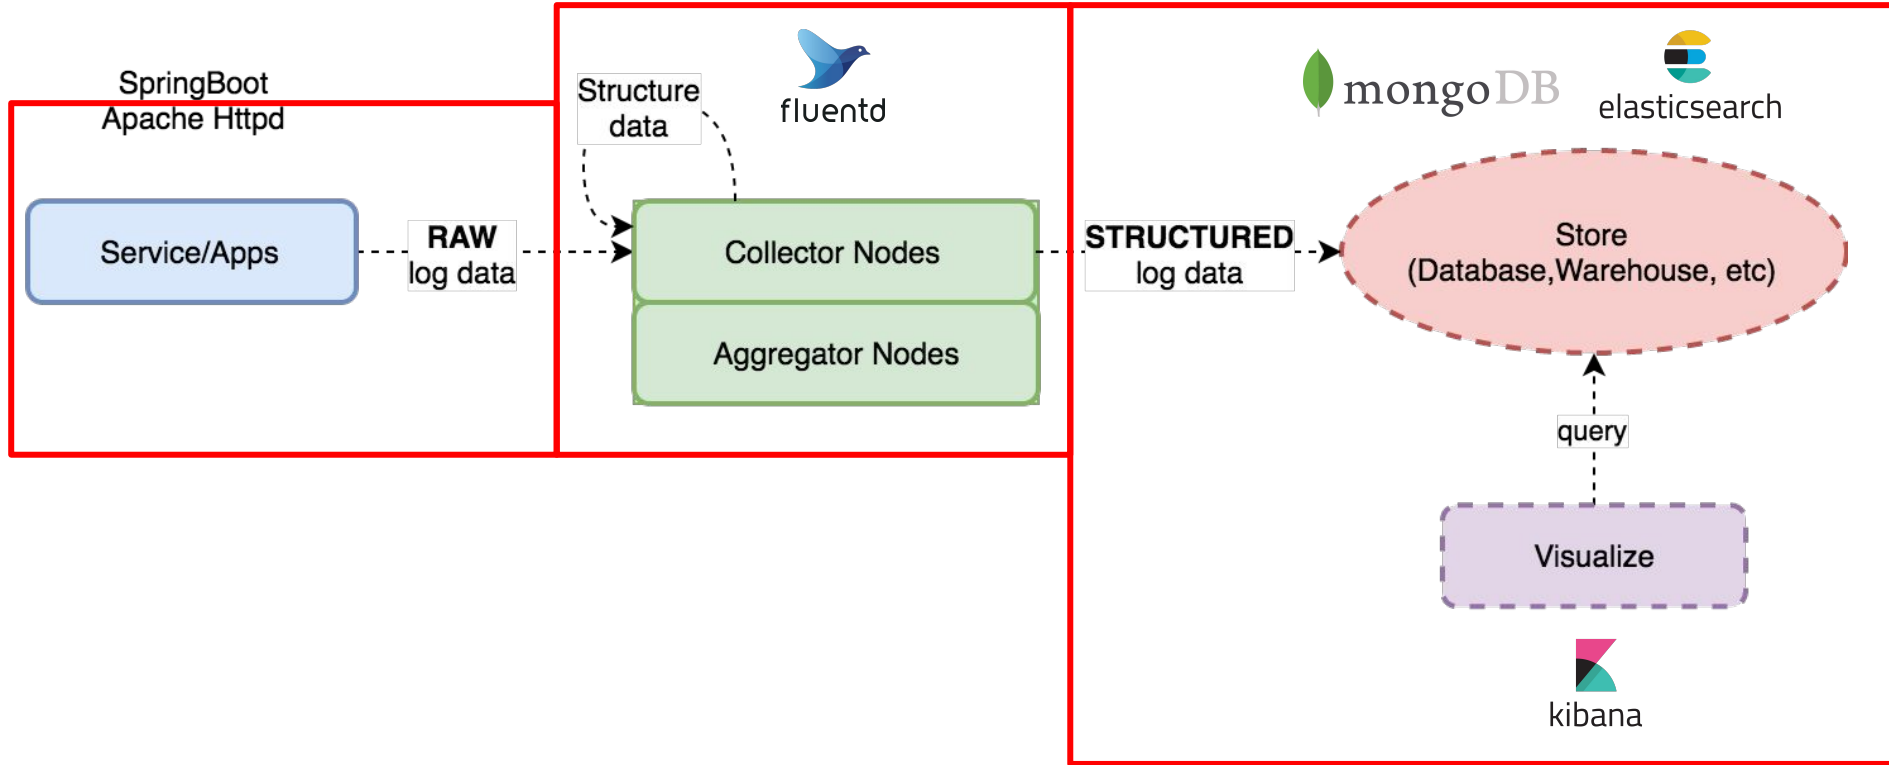
```
    include_time_key true
```

```
  </store>
```

```
</match>
```

Demo

Demo: Capture HTTP Access Logs -> ELK stack



```
# file: fluent.conf
# code intentionally omitted
```

```
<match http.access>
```

```
  @type copy                                → Copy to multiple destinations
```

```
  <store>
```

```
    @type stdout                            → Console output
```

```
  </store>
```

```
  <store>
```

```
    @type elasticsearch                    → Elasticsearch output
```

```
    host elasticsearch
```

```
    port 9200
```

```
    flush_interval 5
```

```
    logstash_format true
```

```
    include_tag_key true
```

```
  </store>
```

```
  <store>
```

```
    @type file
```

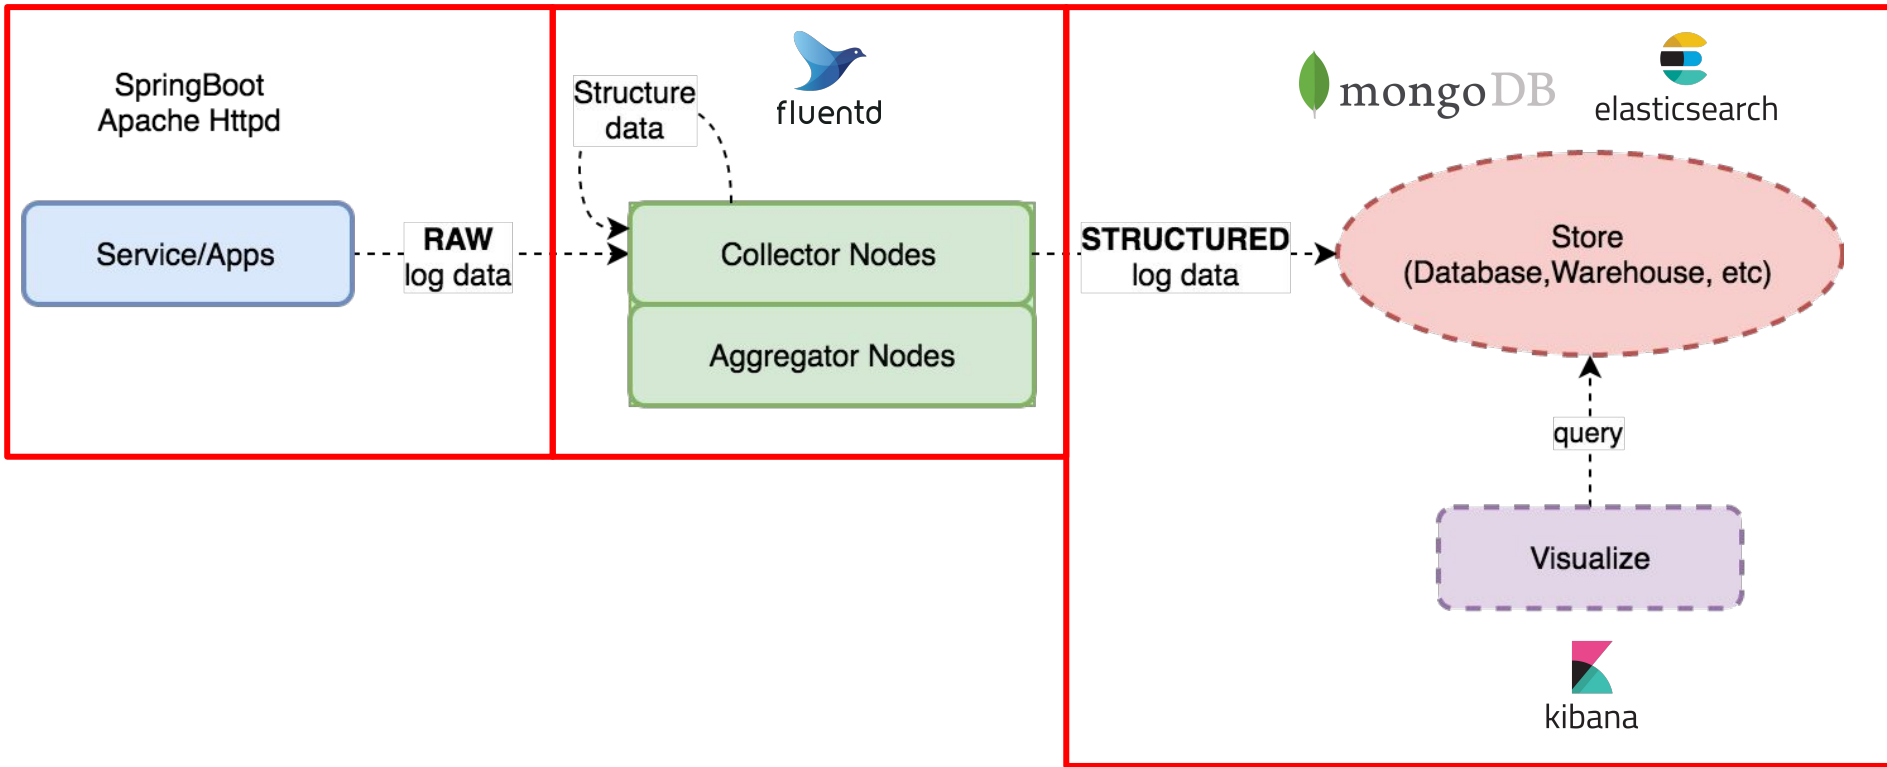
```
    path /fluentd/etc/logs/                → File output
```

```
  </store>
```

```
</match>
```

Demo

Demo: Capture Spring Boot Logs




```
# file: fluent.conf
# code intentionally omitted
```

```
<filter springboot.**>
```

```
  @type parser
```

```
  key_name log
```

```
  reserve_data true
```

```
  reserve_time true
```

```
  <parse>
```

```
    @type grok
```

```
    grok_failure_key grokfailure
```

```
    <grok> → Parsing done based on GROK Patterns
```

```
      pattern %{TIMESTAMP_ISO8601:time_stamp}%{SPACE}%{LOGLEVEL:log_level}...*)
```


```
    </grok>
```

```
  </parse>
```

```
</filter>
```

Demo

Demo: HA Setup



Demo

That's a wrap!

Question?

Marco Pas

Philips Lighting

Software geek, hands on
Developer/Architect/DevOps Engineer

@marcopas